

Boot Flow for P8

v1.52 (11/13/14)

This document is for OpenPower use only.

Copyright and Disclaimer

© Copyright International Business Machines Corporation 2014

IBM, the IBM logo, and ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the Web at “Copyright and trademark information” at www.ibm.com/legal/copytrade.shtml.

Linux is a registered trademark of Linus Torvalds in the United States, other countries, or both.

Other company, product, and service names may be trademarks or service marks of others. All information contained in this document is subject to change without notice. The products described in this document are NOT intended for use in applications such as implantation, life support, or other hazardous uses where malfunction could result in death, bodily injury, or catastrophic property damage. The information contained in this document does not affect or change IBM product specifications or warranties. Nothing in this document shall operate as an express or implied license or indemnity under the intellectual property rights of IBM or third parties. All information contained in this document was obtained in specific environments, and is presented as an illustration. The results obtained in other operating environments may vary.

While the information contained herein is believed to be accurate, such information is preliminary, and should not be relied upon for accuracy or completeness, and no representations or warranties of accuracy or completeness are made.

Note: This document contains information on products in the design, sampling and/or initial production phases of development. This information is subject to change without notice. Verify with your IBM field applications engineer that you have the latest version of this document before finalizing a design.

THE INFORMATION CONTAINED IN THIS DOCUMENT IS PROVIDED ON AN “AS IS” BASIS. In no event will IBM be liable for damages arising directly or indirectly from any use of the information contained in this document.

IBM Systems and Technology Group 2070 Route 52, Bldg. 330 Hopewell Junction, NY 12533-6351

The IBM home page can be found at ibm.com®.

1 Introduction

1.1 Description

This document will describe the high-level IPL flow for the IBM POWER servers based on P8 (i.e. PGP chips). It is not intended to contain all low-level details, but instead is designed to illustrate the relationships between various low-level procedures. The source code and specifically the hardware procedures will always be the final documentation.

This document covers both the hardware and firmware flow required to boot a system to the hypervisor state. This includes full energy management capability and enough resources to boot partitions. Note on accuracy : All details down to the procedure call are correct and this document is considered an authoritative reference. Any details within an individual procedure are informational only, the final authority lies within the procedures themselves and their associated reference documents.

This version of the document will cover all POWER8 systems. Please note that this document has a lot of low level details on the initialization of the POWER processor and it's memory subsystem. There are a lot of terms and details in here that are very IBM and POWER centric. We attempted to put as much in the glossary as possible but please feel free to use the mailing list for any questions.

Throughout the document you will see references to a "SP". This stands for a service processor and when used it's applicable to either the FSP (Flexible Service Processor – used within IBM POWER based servers) or the BMC (the OpenPower service processor). For the most part we've tried to remove FSP specific references for the OpenPower work but some may still remain for reference in here.

There are also quite a few references to "Secureboot" throughout this document. Secureboot is a function built into the POWER8 chips to ensure the appropriate firmware is running. Certain aspects of this have been implemented but other portions of it have been saved for another day.

Reading over the Hostboot Programmers guide (same document repo) is recommended prior to reading this document

This version of the document will cover the following systems:

- Palmetto – Open Power Reference Board
- Habanero – OpenPower System

The following diagram gives a high level overview of the IPL flow. The minute details are explained in the rest of the document.

BMC:

BMC High Level Boot Flow

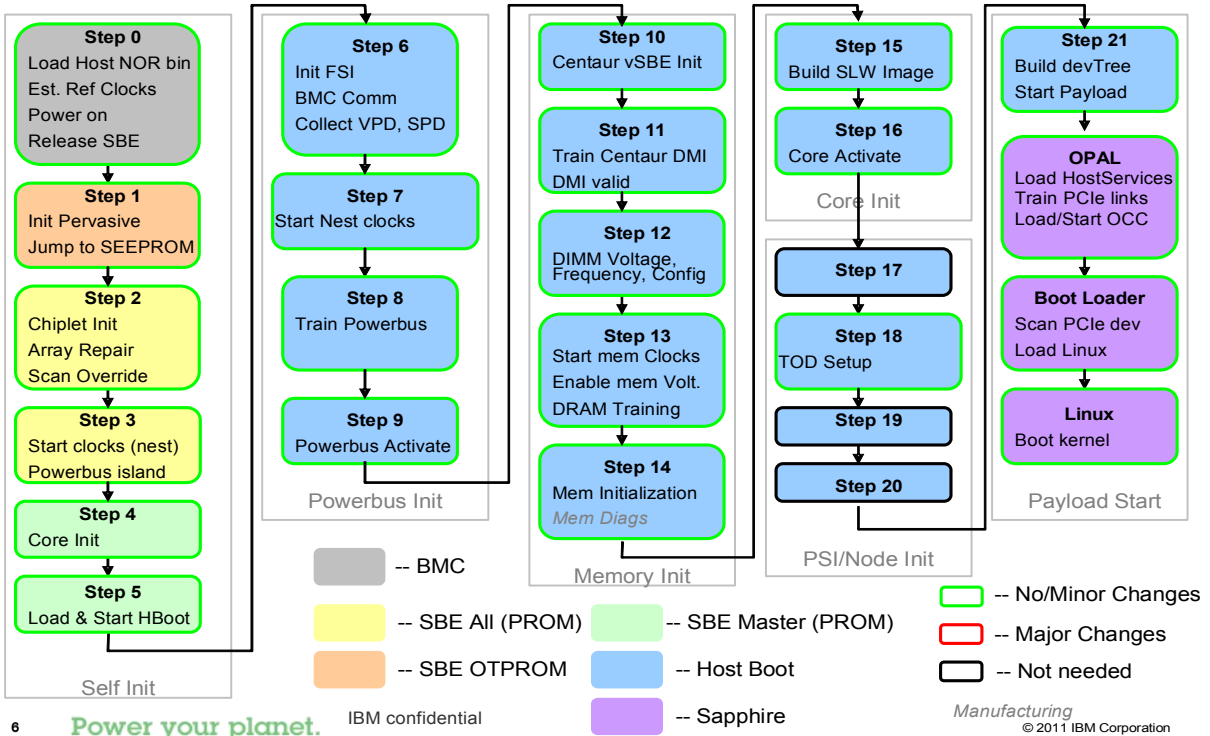


Illustration 1: BMC Boot Flow

1.2 Terminology

- **Centaur:** Memory buffer chip which optimizes memory bandwidth and usage (P8 memory controller connects to it over the DMI bus)
- **Cronus:** Lab debug tool used to control and debug the IPL (used by HW engineers)
- **DMI:** Dynamic Memory Interface bus
- **EX:** Contains the processor core, L2, and L3 logic
- **GPE:** General Purpose Engine – type of PORE unit used by the OCC to offload sequences of scom operations
- **Hardware Reconfig Loop:** when a piece of hardware is deconfigured during the IPL, the firmware can in certain instances go back to a previously run istep and still boot the system with the hardware now deconfigured
- **Hostboot/HB:** FW that initializes the memory and powerbus
- **HB Runtime Services:** Portion of Host Boot that remains resident during hypervisor execution and provides PRD functionality for CEs
- **IPL :** Initial Program Load = Boot process. Covers time between power on and running the hypervisor
- **istep :** IPL Step defined by ecmd interface
- **Maintenance-mode :** IPL that includes extensive diagnostics to test the hardware
- **MPIPL:** Memory-Preserving IPL. This is an IPL type used to grab the hypervisor's memory in certain error scenarios. This is not supported on OpenPower systems.
- **Normal-mode :** IPL that includes minimal diagnostics, focused on functional requirements only

- **OCC**: On Chip Controller – PPC 405 processor that controls the power management per chip
- **PCB**: Pervasive Control Bus – internal processor bus that provides a communication layer between the internal logic within the chip
- **Pegasus** : Family of server products based on around P8 or P8+ processor chips
- **PgP** : P8 Processor Chip
- **PNOR**: P8 Processor NOR chip. NOR flash device where all firmware, including hostboot firmware, is stored and from which it is loaded. It is attached to the master processor through an LPC → SPI bus connection. Called PNOR to distinguish from other NOR chips in the system
- **PORE**: Power On Reset Engine – Mini-execution engine within each P8 chip used to execute firmware which assists with some of the IPL steps. There are three flavors: SBE, SLW, and GPE that each have different purposes.
- **SBE**: Self Boot Engine – A version of a PORE within each P8 chip which is used to do some basic initialization to each chip and to load and start the Hostboot firmware
- **SLW**: Sleep/Winkle – Refer to the PORE engine that moves the EX units between the running (full power), sleep (core power saved, cache active), and winkled (core and caches powered off) states

1.2.1 IPL Types

IPL Type	SP Power	CEC Standby Power	CEC Logic Power	Mainstore Contents	Applicable Platform
Standby POR	Off -> On	Off -> On	Off	Off	Pegasus
Cold IPL	On	On	Off -> On	Cleared	Pegasus
Warm IPL or Warm Re-IPL	On	On	On	Cleared	Pegasus

1.2.2 Nomenclature/Conventions

- Items in bold-italics represent deliveries from the hardware team, either procedures or engineering data files. Ex: ***p8_cfaminit.C***
- Initfile processing is shown with the “WHEN=” value in parentheses, e.g. galaxy2.initfile(L) means to process the initfile with “WHEN=L”.
- Steps in italics are software only and do not interact with the hardware at all.
- istep commands in blue are performed by the Self Boot Engine (SBE)
- istep commands in green are performed by the host code.
 - Note for multi-drawer or multi-Hostboot instances all commands are issued in parallel to all instances, except where otherwise noted
- istep commands in black are performed by the attached service processor

2 Standby POR

BMC Based Systems

This flow starts as soon as the power supplies are plugged in and ends when the BMC is ready to power on the host

1. Apply AC Power

Boot Flow for P8

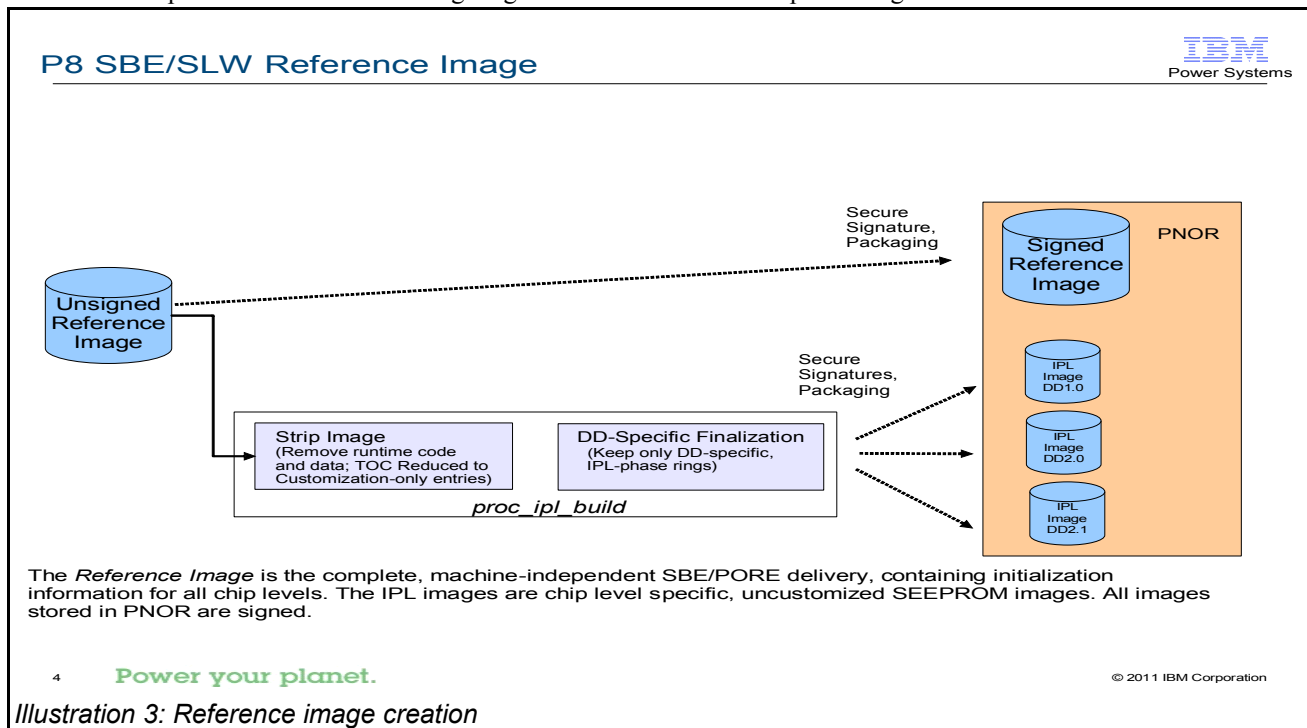
- a) Standby regulators power on from 12VCS
- b) Reset generator starts 200ms after Standby Pgoods high
- c) Standby reset feeds into BMC/APPS. (any intelligent device). BMC/APSS running.
- d) P8 and centaur CFAMS are NOT powered
2. APSS loads itself from internal flash.
3. BMC starts running from SPI Flash at address 0.
4. BMC loads u-boot
5. U-boot loads the Linux kernel
6. AMI framework loads. This includes the following daemons:
 - a) IPMI
 - b) LPC bus access setup for Host to access PNOR. This must be dynamic based on size of PNOR
 - c) PNOR mounted via mtdblock
7. System now at Standby

At the end of this flow the SP and DPSS chip are powered on and viable. For the purposes of this discussion SP Standby is just enough information to power on the CEC logic. SP has **not** issued the “GO” bit yet.

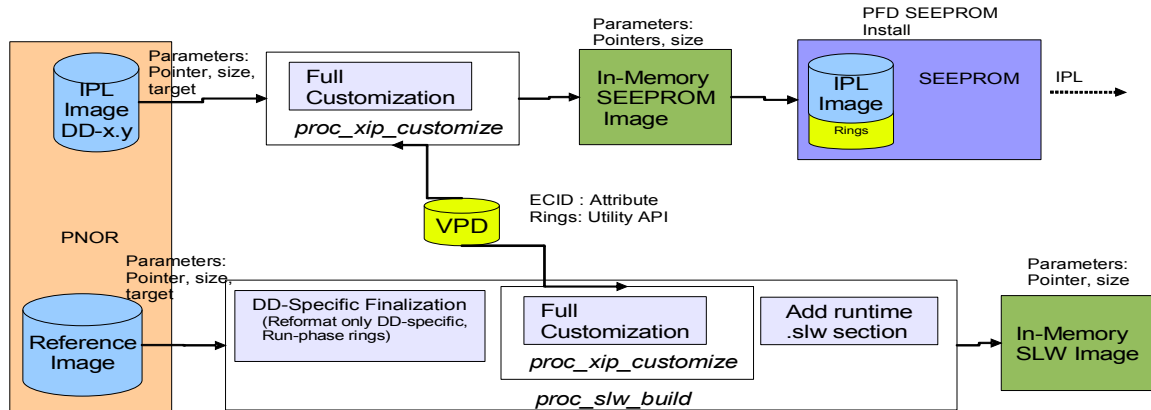
If this is a Pegasus system the PgP and Centaur chip's CFAM portion is powered on and has FSI clocks.

3 EX Initialization Image

The EX's are initialized via the Sleep Winkle (SLW) Engine from the SLW reference image. The image is setup by Hostboot in istep 15 and 16. The following diagrams show the relationship and usage.



Reference and SEEPROM Image customization



proc_xip_customize is run by hostboot and FSP (at the factory) to customize an in-memory SBE-XIP format binary, which is then written to SEEPROM. *proc_slw_build* is run by host code to convert the reference image into a DD-specific in-memory SLW image.

5

Power your planet.

© 2011 IBM Corporation

Illustration 4: Image customization

4 Cold IPL

BMC Based

This flow covers the steps from BMC Standby through the initial load of OPAL

0. Power on

- If PNOR is connected via BMC buses, the BMC must ensure that PNOR is accessible via FW LPC operations starting at 0xFFFF_FFFF and downward. The Hostboot Base image must be at address ECC adjusted address 0xFFF78000, which is LPC address 0xFFF67000, which is SPI address 0x03f67000 for a 64MB NOR or 0x01f67000 for a 32MB NOR.
- BMC raises all CEC power rails except for VDDR.
- BMC issues FSI Magic start sequence to the master P8 chip
 - The master P8 chip is the one connected via LPC to the BMC/PNOR
 - The start sequence consists of selecting the SBE boot side and then triggering the SBE boot

1. Self Boot Engine – OTPROM

- Initialize pervasive logic
- Reset I2C bus to SEEPROM
- Jump to execution from SEEPROM

2. Self Boot Engine – SEEPROM

- Chiplet initialization
 - Flush to 0s

- Array repair
- Init fixes

3. Self Boot Engine – SEEPROM

- a) Start clocks (nest)
- b) Chiplet enters powerbus island mode

4. Self Boot Engine – EX Init

- a) Only runs on master chip
- b) Pick EX as master EX and run through SLW initialization

5. Self Boot Engine – Hostboot load

- a) Only runs on master chip
- b) Assumes that the BMC has enabled PNOR controller prior to releasing the SBE:
 - Configure the direct read windows on LPC address spaces
- c) Enable the LPC interface to permit access to the PNOR
 - Releases LPC reset
- d) SBE fetches payload (Hostboot base image) from PNOR flash and stuffs it into targeted EX's L3
 - SBE payload must have a wrapper that contains a small header containing image size followed by target real address. This is above and beyond the Secureboot wrapper
 - SBE header is as follows:
 - 8 bytes – unsecure mem address
 - 8 bytes – secure mem address
 - 8 bytes – image size
 - Goal of the SBE header is to allow a common SBE code regardless of Secureboot enablement. The SBE will use the secure scom register (0x00010005) to determine if Secureboot is enabled or not (based on physical jumper). For example these are the values when loading an image to 128MB. All images **must** have the 4K secureboot header, even for unsecure images – even if it is just stubbed
 - Unsecure mem address – 0x 180000 (1.5MB)
 - Secure mem address – 0x8000000 (128MB)
 - size – 0x81000 (512KB + 4KB for secureboot header)
 - Secure boot off (Mimic ROM):
 - Core Scratch Reg 0 – Xscom base address
 - Core Scratch Reg 1 – Undefined (not used)
 - HRMOR – Secure load address
 - Write Hostboot image – Secure mem address -4KB
 - Secureboot on(load to unsecure, ROM moves to proper location)
 - Core Scratch Reg 0 – Xscom base address
 - Core Scratch Reg 1 – Unsecure load address
 - HRMOR – Undefined (not used, set by ROM)
 - Write Hostboot image – Unsecure mem address - 4KB

- Image is ECC protected
 - HB Base Image is core of Hostboot – enough code to get off the ground
 - Includes base kernel/OS, Xscom driver, PNOR driver, utilites
 - Hardware procedures are embedded in extended image. SBE does **not** load this
 - Updating procedures will require generation of image and flash write to PNOR
 - e) Instruction start on one core, one thread. Once hostboot gets control of the thread (either from Secureboot ROM or immediately) it will load and then start the other threads on the core
 - f) After SBE starts instructions on the core it execute the master winkle assist procedure in the background
 - The procedure will “hang” in a loop waiting for a trigger from Hostboot to move onto its active logic
 - See istep 15 for more detail
6. **Step 6 Hostboot – Slave SBE**
- 6.1. `host_setup (non executable istep): Setup host environment`
- a) If in secure boot the ROM as already validated image
 - b) Select primary thread (only thread running)
 - c) Initial setup
 - stacks
 - MSR
 - execution environment
 - Thread control structures
 - Memory Management setup
 - Console (BMC only)
 - d) Ready for execution
 - Start and release all other threads on core (1-7)
 - Tracing
 - Device Drivers
 - Xscom
 - Mailbox
 - e) HB mechanism to read/write to PNOR
 - All Xscom based
 - Host writes to LPC ↔ SPI NOR controller to read/write
 - SBE uses NOR at lowest frequency, Hostboot will use flash config info to speedup to full frequency
- 6.2. `host_istep_enable (non executable istep): Hostboot istep ready`
- a) Hostboot checks PNOR for istep attribute, if set Hostboot “halts” and waits for commands from SP
 - istep mode can also be entered by setting the appropriate bit in the SIO scratch regs on BMC based systems
 - b) Only isteps after this point can be issued to Hostboot
 - c) At this point communication can be performed with the SP
 - For FSP this is via the CFAM FSI2Host mailbox
 - For BMC this is via the Block Transfer IPMI protocol
- 6.3. `host_init_fsi` : Setup the FSI links to slave chips
- a) It is expected that the following steps have already been done by FSP – Hostboot will just use FSI bus.

For BMC based systems the Hostboot will perform these actions

- Configure FSI master (HUB and Cascade)
 - Send break commands to FSI slaves
 - Configure the slaves
 - Force lbus
- b) Setup Scm device drivers
- Read ID/EC levels
- c) Reset all I2C engines/slaves on the P8 Master Chip and all FSI I2C Masters (P8 slaves, centaurs)
- Can't reset the scm only I2C master on the P8 Slave chips (see 6.12)
- 6.4. `host_set_ipl_parms` : Build ipl parameters
- a) Sets the IPL parameters for this boot
- 6.5. `host_discover_targets` : Builds targeting
- a) Determines what targets are present and functional
- b) This is the step where the host “configures” itself and builds its present/functional map of the targets
- Uses FSI presence to detect processors and memory buffers
 - For OpenPower systems Hostboot reads the VPD and updates its cache in PNOR if the part/serial number does not match. All VPD is read across the FSI bus except for the master processor, which is done via Xscom
 - Processors – MVPD
 - Centaurs – board VPD
 - Dimms – SPD
 - Compares data in PNOR with FSI presence detect. If there is a mismatch then an error is flagged and part is deconfigured
- c) For OpenPower systems Hostboot will push the IPMI FRU inventory to the BMC
- Must push for all present parts
 - Must update FRU present/functional state
- 6.6. `host_gard` : Do Gard
- a) Apply repeat-gard records and deconfigure hardware
- 6.7. `host_cancontinue_clear` : Clear deconfigured states
- a) Clears Cancontinue target deconfigured state (used to know when to trigger a CanContinue IPL)
- 6.8. `proc_revert_sbe_mcs_setup` : Clean up MCS extent regs
- a) *`proc_revert_sbe_mcs_setup.C`*
- Clean up the MCS BARs that were used by SBE and Hostboot to cleanly load/purge the L3 cache
 - Re-enable speculative reads
- 6.9. `proc_cen_ref_clk_enable` : Setup centaur ref clocks
- a) *`proc_cen_ref_clk_enable.C`*
- Enable the ref clocks to centaur
- 6.10. `host_slave_sbe_config`
- a) Need to run this from master processor to all slave processors for Secureboot hole (need to ensure that FSP didn't leave compromised P8 Slave. Also run on all BMC systems to start the slave processors since the BMC only starts the master
- b) *`proc_setup_sbe_config.C`*
- Update SBE config data area with any configs/parameters required by SBE (see step 0 for more

details)

6.11. `host_sbe_start`

a) `proc_start_sbe.C`

- Set a bit to start the SBE engine on slave chips. Located in FSI GP region
- This same bit performs the scan0 flush of pervasive

6.12. `proc_check_slave_sbe_seeprom_complete` : Check Slave SBE Complete

a) `proc_check_slave_sbe_seeprom_complete.C`

- Check to make sure that the slave SBE engines have completed
 - SBE state reg to make sure that slave SBE has finished
 - Check for specific “done” signature in status reg (not tied to istep number)
- If they haven't wait 1 second and then generate error
- Also checks for error status codes

b) `proc_extract_sbe_rc.C -soft_err`

- Called on all chips (master and slaves) to look for any correctable errors on the PNOR and/or SEEPROM
- The `soft_error` flag just tells the procedure to not generate an error if no HW issue

c) Reset all scom only I2C engines/slaves on the P8 Slave Chips

6.13. `proc_xmit_sbe` : vSBE Init of Slave Chips

a) Left as placeholder if step 3 is needed by master

b) `proc_xmit_sbe.C`

- Loads SBE PNOR image from PNOR into memory (L3)
 - Scan escape
 - Customizes parameters for slave chip (system settings, core good, etc)
- Uses virtual vSBE engine to setup slave chips
- vSBE provides callbacks that Hostboot fills in
 - Sends output stream of get/putscoms to slave chips across FSI to slave chips
 -
 -

7. Step 7 Hostboot – Nest Chiplets

7.1. `proc_attr_update` : Proc ATTR Update

a) `proc_attr_update.C`

- Called per processor
- Stub HWP for FW to override attributes programatically

7.2. `proc_a_x_pci_dmi_pll_initf` : PLL Initfile for A, PCIe, DMI

a) `proc_a_x_pci_dmi_pll_initf.C`

- PLL rings are stored as system attributes
 - Included tune bits, frequency

7.3. `proc_a_x_pci_dmi_pll_setup` : Setup PLL for A, X, PCIe, DMI

a) `proc_a_x_pci_dmi_pll_setup.C`

- Checks that the PLL locked
- Start the VAR OSCs / Config the TANK PLLs & lock

- In certain configs these chiplets are potentially not used
 - Must run at system frequency
- 7.4. `proc_startclock_chiplets` : Start clocks on A, X, PCIe chiplets
- a) `proc_startclock_chiplets.C`
- Start Xbus, Abus, PCIe clocks
 - Start clocks on configured chiplets for all chips (master and slaves)
 - Drop fences (RI/DI)
- 7.5. `proc_chiplet_scominit` : Apply scom inits to chiplets
- a) `proc_chiplet_scominit.C`
- Initfiles in procedure defined on VBU ENGD wiki ([TODO add link](#))
 - Apply scom overrides to all good chiplets (except EX and MC)
 - Powerbus – Apply system Pbus config to slave PB (scoms to CD NEXT)?
- b) `proc_psi_scominit.C`
- Each instance of bus must have unique id set for it – personalize it
 - Must set present and valid bits based on topology (Attributes indicate present and valid)
- 7.6. `proc_xbus_scominit` : Apply scom inits to Xbus
- a) `proc_xbus_scominit.C`
- Each instance of bus must have unique id set for it – personalize it
 - Must set present and valid bits based on topology (Attributes indicate present and valid)
- 7.7. `proc_abus_scominit` : Apply scom inits to Abus
- a) `proc_abus_scominit.C`
- Each instance of bus must have unique id set for it – personalize it
 - Must set present and valid bits based on topology (Attributes indicate present and valid)
- 7.8. `proc_pcie_scominit` : Apply scom inits to PCIechiplets
- a) `proc_pcie_scominit.C`
- Initfiles in procedure defined on VBU ENGD wiki ([TODO add link](#))
 - Perform the PCIe Phase 1 Inits 1-8
 - Sets the lane config based on MRW attributes
 - Sets the swap bits based on MRW attributes
 - Sets valid PHBs, remove from reset
 - Performs any needed overrides (should flush correctly) – this is where initfile may be used
 - Set the IOP program complete bit
 - This is where the dSMP versus PCIE is selected in the PHY Link Layer
- 7.9. `proc_scomoverride_chiplets` : Apply sequenced scom inits
- a) `proc_scomoverride_chiplets.C`
- Apply any sequence driven scom overrides to chiplets – Should be NONE
8. **Step 8 Hostboot – EDI, EI Initialization**
- 8.1. `fabric_erepair` : Restore Fabric/EDI Bus eRepair data
- a) Restore/preset bad lanes on A, and X, buses from VPD
- `io_restore_erepairr.C` (A, X bus target pairs)
 - procedure for repairs on X and A buses

- Applies powerbus repair data from module vpd (#ER keyword in VRML VWML)
 - Applies centaur data from planar prom (planar centuars), centuar dimm prom (centuar dimm)
 - Runtime detected fails that were written to VPD are restored here
- 8.2. `fabric_io_dccal` : Calibrate Fabric/EDI interfaces
- a) `io_dccal.C` (A, X bus target pairs passed in)
- Will be called per bus target pair
 - For CCM the A bus `io_dccal` must be done by the SP for “hot add/remove” buses on active nodes. HB does new node side
 - Calibration of TX impedance, RX offset for A and X busses
 - Needs to be quiet on the bus – drivers are quiesced and driving 0s – A, X buses
 - Must be complete on ALL chips before starting next A, X bus training
 - Expect to use a calculation
 - At end of offset calibration there may be a lane that is bad
 - FW must record bad lane and write to VPD for future eRepair (handled when PRD starts)
 - Must generate error log, procedure will mark lane bad in HW (which future procedure take advantage of)
- b) `io_new_procedure.C?`
- 8.3. `fabric_pre_trainadv` : Advanced pre EI/EDI training
- a) `io_pre_trainadv.C` (called on each A and X bus target pair)
- Debug routine for IO Characterization
 - Nothing in it
- 8.4. `fabric_io_run_training` : Run training on internal buses
- a) `io_run_training.C` (called on each A and X bus target pair)
- **Hostboot will run training on all intra node buses. For OpenPower this is all A buses. It can include X buses as well. For multinode systems this is run by the FSP in a later step**
 - Start wder on all slave bus endpoints
 - Start wder on all master bus endpoints
 - Wiretest, Deskew, Eye Optimization, and repair
 - Option to run extend bit patterns in optimization phase (replaces RDT)
 - Repairable fails are left for PRD to analyze and move data into VPD
 - PRD will use `io_eRepair_read.C` to perform this
 - Fatal bus training errors are handled by procedure, must return error and FFDC (written to VPD for PgP)
 - Expected that fatal error passes returncode back to HWPF, FW then looks up returncode and determines what to do based off of errorInfoRepository
 - **TODO FW team: Make provisions to handle some sort of retry that has to reset both sides**
- 8.5. `fabric_post_trainadv` : Advanced post EI/EDI training
- a) `io_post_trainadv.C` (called on each A and X bus target pair)
- Debug routine for IO Characterization
 - Nothing in it
- 8.6. `host_startprd_pbus` : Load PRD for powerbus domain
- a) Hostboot will apply fabric topology configuration based on Abus, Xbus, and processor chip deconfiguration

- b) Start Host PRD just for powerbus domain
- c) Must be minimal environment (running out of L3)
- 8.7. `host_attnlisten_proc` : Start listening for attentions
 - a) Enable hostboot to start including processor (all) attentions in its post istep analysis
 - b) From this point on ATTN/PRD will listen (“poll”) for powerbus attentions after each named istep
- 8.8. `proc_fab_iovalid` : Lower functional fences on local SMP
 - a) `proc_fab_iovalid.C` (called on A and X bus target pair)
 - Lower fences on internal fabric X buses
 - Lower fences on external fabric (A) buses
 - Only performed on trained, valid buses
 - b) After this point a checkstop on a slave will checkstop master

9. Step 9 Hostboot – Activate PowerBus

9.1. `proc_build_smp` : Integrate PgP Islands into SMP

- a) `proc_build_smp.C`
 - Look for checkstops
 - CCM all PgP islands into the SMP
 - Fabric config between dSMP and IO/CAPI are set here – only can set once, must be known by this point in time
 - After this point the SMP is built for normal mode
 - High level flow
 - Preconditions:
 - Slave powerbus init'ed via vSBE
 - Slave chips Pbus have run fabric init via vSBE
 - CD CURR/NEXT is “island config” and AB CURR/NEXT is “island” mode
 - EI/EDI buses are up and alive
 - HB is loaded on master with AB CURR as island mode
 - Apply system Pbus config to both master/slave PB (scoms to CD NEXT)
 - Use ADU to switch CD to NEXT on **master chip** (Master Pbus has system config)
 - via FSI use ADU to switch CD to NEXT on all **slave chips** (Slave chip have system config)
 - Via FSI quiesce slave fabric
 - Via FSI apply system fabric topology to AB NEXT and AB CURR setting on slave chips
 - Via Xscom apply system fabric topology to AB Next on master
 - Apply switch AB to master via ADU
 - Full powerbus topology comes alive following init in the switch AB
 - Via Xscom write master AB NEXT to system fabric topology

9.2. `host_slave_sbe_update`

- a) Hostboot must update SEEPROM because the SP cannot because of secureboot. It is at this step in the IPL so it can be updated via Xscom (trusted path) on all chips in the system
- b) `p8_customize_image.C`
 - If needed build a custom SEEPROM image for each chip in the system off of the base IPL SEEPROM image
- c) If the SEEPROM was updated then Hostboot will request a reipl at this point
- d) On systems that support Alt Master processors then code will attempt to read the TOC of the Alt Master PNOR to check for connection problems. If an error is detected it will be logged, but this does not stop

the IPL (except when in manufacturing mode)

e) *proc_pcie_hotplug_control.C – off*

- Called on all functional processors to turn off PCIe hotplug controllers. Done to workaround PERST spec violation in P8 chip. Uses scom logic to drive the host I2C master to power off the hotplug controllers. May be done as I2C DD commands as well.

10. Step 10 **Hostboot SBE – Centaur Init**

The following steps are part of the Centaur SBE Init Image. Note that these cannot be stepped in hostboot, but can be in Cronus (due to POREVE implementation). When running hostboot when it is requested to execute step 10.2 it will perform all substeps in step 10. From a command line interface all the rest of the steps should be stubbed and return success. Only when running in Cronus only (vSBE) mode will steps 10.3-10.x do anything. The HW reconfig loop is only possible on processor DD2.x or higher parts. FW is responsible for disabling the HW reconfig loop on D1.X parts.

Hostboot will check for HW reconfig loop after then end of each named istep. If it happens in 10,11,13, or 14 it will go back to the beginning of step 10 (known as HW reconfig loop). If it fails in step 12 it will go back to the beginning of step 12 (FW reconfig loop).

Note that this is the step for HW Reconfig to restart on centaur or dimm training/init fails

The following sections are to initialize the Centaur TP logic (Host TP Initialization)

10.1. *host_prd_hwreconfig* : Hook to handle HW reconfig

- a) This step is always called
- b) Move all centaur's inband scom back to FSI scom
- c) Call PRD to allow them to rebuild model to remove centaurs
- d) Used for HW reconfig path. FW's strategy is to perform the reconfig on ALL functional centaurs/MCS's in the system. The following procedures must be called:
 - *proc_enable_reconfig.C*
 - Call on all present processor, MCS, and centaur targets. Called per processor and associated targets
 - Enables HW for reconfig loop
 - Add a Chip EC feature attribute to check support of special MCS bit
 - Check this when setting the bit to bring the DMI bus down
 - Add an attribute (ATTR_MSS_INIT_STATE) to each centaur to track where the Reconfig loop got too:
 - Clocks on (can do fir masking) – set after step 10
 - DMI bus up (inject special bit) – set after framelock
 - Turn's on special bit that allows the MCS DMI to get errors and not get into a hang condition
 - Mask a bunch of FIRs on processor
 - Mask a bunch FIRs on centaur (HWP will check clock state)
 - Injects a fail on the DMI bus (only if DMI bus is alive)
 - Clears IO/MCS FIRs
 - Turns off special bit

cen_sbe_istep_pnor.S

- This procedure is not an official istep under FW/Cronus control
 - It is a wrapper that allows for individual istep control
 -
- 10.2. *cen_sbe_tp_chiplet_init1* : TP Chiplet Init
- a) *cen_sbe_tp_chiplet_init1.S*
- Flush all GP reg content to default state
- 10.3. *cen_sbe_pll_initf* : Program Nest PLL
- a) *cen_sbe_pll_initf.S*
- Apply the Nest PLL ring
 - Initfile defined on VBU ENGD wiki ([TODO add link](#))
 -
- 10.4. *cen_sbe_pll_setup* : Setup Nest PLL
- a) *cen_sbe_pll_setup.S*
- Performs PLL checking
 - The memory PLL (ie DDR3) are set to a default value (not actual runtime speeds)
 - Establish Nest PLLs (feeds TP chiplet)
 - Includes scans for all PLL (Nest, EDI) setup
 - Moved out of bypass(just nest)
- 10.5. *cen_sbe_tp_chiplet_init2* : Cen TP Chiplet Init 2
- a) *cen_sbe_tp_chiplet_init1.S*
- Scan 0 init TP unit (flush)
 - No repair/timing for TP chiplet (ie fuses)
- 10.6. *cen_sbe_tp_arrayinit* : Cen TP Chiplet array init
- a) *cen_sbe_tp_arrayinit.S*
- Run arrayinit on TP chiplet
 - After this all arrays are initialized
 - Scan flush 0 to all rings except GPTR, Time, Repair
- 10.7. *cen_sbe_tp_chiplet_init3* : Cen TP Chiplet Start clocks
- a) *cen_sbe_tp_chiplet_init3.S*
- Start clocks on perv region
 - At end, the TP chiplet can be used to init the rest of the chip
 - Involves writing FSI GP3 to switch mux – all access now go through TP chiplet

The following sections are to initialize the Centaur chip logic (Host Chiplet Setup)

- 10.8. *cen_sbe_chiplet_init* : Cen Chiplet Init
- a) *cen_sbe_chiplet_init.S*
- Scan 0 all rings on all good chiplets (except for TP)
 - Kick off the fuse repair algorithm (loads repair ring)
- 10.9. *cen_sbe_arrayinit* : Chiplet array init
- a) *cen_sbe_arrayinit.S*
- Run arrayinit on all good chiplets (except for TP)

- After this all arrays are initialized
 - Scan flush 0 to all rings except GPTR, Time, Repair
 - *Note the MBA PLLs are controlled later by Hostboot*
- b) *Note: If LBIST was to be run, it should be run after this step, prior to the next step*

The following sections are to initialize the Centaur chip logic (Host Chiplet Initialization)

10.10. *cen_sbe_dts_init* : *Cen DTS Init*

a) *cen_sbe_dts_init.S*

- Setup the thermal sensors based on fuse data

10.11. *cen_sbe_initf* : *Cen Scan overrides*

a) *cen_sbe_chiplet_initf.S*

- Perform any scan overrides for Centaur
 - May not have any config dependent scans
- Does not include the pervasive region

10.12. *cen_sbe_do_manual_inits* : *Manual Cen Scans*

a) *cen_sbe_do_manual_inits.S*

- Perform any non initfile scan overrides for Centaur
- Should be avoided, provided as a placeholder for workarounds

The following sections are to initialize the Centaur chip Scm logic

10.13. *cen_sbe_nest_startclocks* : *Start Cen Nest*

a) *cen_sbe_nest_startclocks.S*

- Starts Centaur's nest clocks. This includes the L4, DMI, nest portions of logic. It specifically excludes the MBA clocks
- Deassert the memrst_b GP bit to activate the reset_OE signal
- Enable driver and receivers (set appropriate GP bits)
- Lower RI and DI inhibits

10.14. *cen_sbe_scominits* : *Perform any Cen scom inits*

a) *cen_sbe_scominits.S*

- Any needed scom initializations – no config dependent settings allowed
- Should be empty
-

11. Step 11 Hostboot – DMI Training

11.1. *mss_getecid* : *Read out ECID of all Centaurs*

a) *mss_get_cen_ecid.C*

- Read the ECID for each centaur and store away for callouts. Also update functional state of MBAs and L3 cache halves for later use
- Sets ATTR_MSS_INIT_STATE to “clocks on”

11.2. *dmi_attr_update* : *DMI ATTR Update*

a) *dmi_attr_update.C*

- Called per MCS and Centaur
 - Stub HWP for FW to override attributes programatically
- 11.3. `proc_dmi_scominit` : DMI Scm setup on P8 MCS
- a) `proc_dmi_scominit.C`
- Initfiles in procedure defined on VBU ENGD wiki ([TODO add link](#))
 - Called per P8 MCS
 - Perform scom inits for DMI on the processor
- 11.4. `dmi_scominit` : Scm setup on centaur
- a) `cen_dmi_scominit.C`
- Initfiles in procedure defined on VBU ENGD wiki ([TODO add link](#))
 - Perform scom inits for DMI on Centaur
- 11.5. `dmi_erepair` : Restore EDI Bus eRepair data
- a) Restore/preset bad lanes on MC buses from VPD
- Bad lanes are preset on the receive side
 - `io_restore_erepair.C (MC and Centaur target pair passed in)`
 - procedure for repairs on MC bus
 - Applies centaur data from planar prom (planar centuars), centuar dimm prom (centuar dimm)
 - Runtime detected fails that were written to VPD are restored here
- 11.6. `dmi_io_dccal` : Calibrate DMI interfaces
- a) `io_dccal.C (MC and Centaur target pair passed in)`
- Calibration of TX impedance, RX offset for memory buses
 - Needed for EDI buses on p8 and centaur (need to run in order to guarantee clean clock)
 - Needs to be quiet on the bus – drivers are quiesced and driving 0s – EDI buses
 - Must be complete on ALL centaurs for this PgP island before starting next EDI training(host based sync point)
 - At end of offset calibration there may be a lane that is bad
 - FW must record bad lane and write to VPD for future eRepair (handled when PRD starts)
- 11.7. `dmi_pre_trainadv` : Advanced pre DMI training
- a) `io_pre_trainadv.C (MC and Centaur target pair passed in)`
- Debug routine for IO Characterization
 - Nothing in it
- 11.8. `dmi_io_run_training` : Run training on MC buses
- a) `io_run_training.C (MC and Centaur target pair passed in)`
- Train internal DMI bus
 - Wiretest, Deskew, Eye Optimization, and repair
 - Option to run extend bit patterns in optimization phase (replaces RDT)
 - Wiretest fails are left for PRD to analyze and move data into VPD
 - Fatal bus training errors are handled by procedure and written to VPD
- 11.9. `dmi_post_trainadv` : Advanced post DMI training
- a) `io_post_trainadv.C (MC and Centaur target pair passed in)`
- Debug routine for IO Characterization

- Nothing in it
- 11.10. `proc_cen_framelock` : Initialize EDI Frame
 - a) *proc_cen_framelock.C*
 - Raise IO Valid – Allow link init traffic (scrambled patterns) on EDI bus
 - P8 Centaur initial frame lock
 - Starts listening automatically after IOValid raised
 - Started on the P8 logic
 - If a bit error (CRC) in the middle need to reFrameLock
 - Round trip delay calculation
 - Host code can trigger and check
 - Inband accesses are now viable
 - Hardware xmitting idle frames
 - Enabled CRC checking
 - EDI is at runtime state
 - If successful, set ATTR_MSS_INIT_STATE to DMI active on centaur
- 11.11. `host_startprd_dmi` : Load PRD for DMI domain
 - a) Expand Host PRD to include DMI (as well as powerbus)
 - b) Must be minimal environment (running out of L3)
- 11.12. `host_attnlisten_cen` : Start listening for attentions
 - a) Enable hostboot to start including centaur attentions in its post istep analysis
- 11.13. `cen_set_inband_addr` : Set the Inband base addresses
 - a) *proc_cen_set_inband_addr.C*
 - Any initializations to setup the Inband access path
 - MCS – Scm base address
 - Centaur – any other settings
 - **ALL ACCESES from this point on in are Inband access for Centaur unless otherwise specified**

12. Step 12 Hostboot – MC Config

Note that the “FW Reconfig” loop starts here (since it doesn't touch HW). Any reconfig during step 12 will loop back to this step

- 12.1. `host_collect_dimm_spd` : Collect Master dimm SPD
 - a) The SPD is not collected here, but the step's name remains for historical reasons
 - On BMC systems the SPD was collected in discover_targets
 - On FSP systems, the FSP has placed all dimm SPD into PNOR prior to releasing the SBE, so Hostboot consumes the dimm SPD directly from PNOR
 - b) Hostboot knows the relationship between dimm, SPD, chip, voltage rail, and socket
 - c) Hostboot FW then places dimm SPD of interest into HWPf attributes
 - This may be “on demand” as attributes are looked up
 - d) *mss_attribute_cleanup.C*
 - Called on all present centaurs
 - Hook to clean up attributes on reconfig loop (set to known state)
 - e) For the inner reconfig loop

- Call PRD to allow them to rebuild model to remove centaurs
 - ***proc_enable_reconfig.C***
 - FW must only call present, non functional (processor, MCS, and centaur) targets
 - Must not be called on functional parts
 - Does the same thing as above (10.1)
- 12.2. ***mss_volt*** : Calc dimm voltage
- a) ***mss_volt.C***
- Procedure is called all the dimms on a voltage rail
 - Calculate rail Voltage and updates rail system attribute
 - Save settings in variables (saved in framework/cache)
 - Procedure handles checking overrides
- 12.3. ***mss_freq*** : Calc dimm frequency
- a) ***mss_freq.C***
- Procedure is called on each MC in the system
 - Looks at voltage and dimm functionality
 - Calculate per memory controller frequency from attributes – picks the frequency bucket to use
 - Save settings in variables (saved in framework/cache)
 - Procedure handles checking overrides
- 12.4. ***mss_eff_config*** : Determine effective config
- a) ***mss_eff_config.C***
- Called per MC
 - Uses HWPF attributes to determine config (SPD/PNOR overrides was stored previously in step 5)
 - Determine effective settings for settings – determine “final” solution
 - Looking at overrides and applies
 - Extent settings are stored as attributes per target (may be changed later based on training)
- b) ***mss_eff_config_thermal.C***
- Called per MBA
 - Perform thermal calculations for the effective config
- c) In order to ensure optimal placement of memory (ie as contiguous as possible for PHYP) the following sequence must be followed
- ***opt_memmap.C***
 - Called with init = true to zero out all relevant attributes
 - ***mss_eff_grouping.C***
 - Called on each P8 target. Note that all P8 base addresses must be zero
 - This is done the first time to get the available sizes behind all P8s
 - ***opt_memmap.C***
 - Called with init=false
 - All P8 on a physical drawer are passed in to procedure
 - FW written procedure that then walks the P8 mappings and orders the base addresses to form as contiguous a memory map as possible
 - This is a FAPI procedure so it can be shared with Cronus

- *mss_eff_grouping.C*
 - Called on each P8 target. Note that all P8 base addresses are now set optimally
 - This will then slot the maps in the correct places
- *mss_eff_mb_interleave.C*
 - Called on each centaur target.
 - This sets up the MBA interleaving internal to the centaur

12.5. *mss_attr_update* :MSS ATTR Overrides

a) *mss_attr_update.C*

- Called per MC
- Stub HWP for FW to override attributes programatically

13. Step 13 Hostboot – DRAM Training

13.1. *host_disable_vddr* : Disable VDDR on CanContinue loops

a) Power off dram. Must be done here to meet JEDEC spec compliance

- Turned off here to handle Reconfig loop for dimm failure
- Only really issued if VDDR is on
- **On MPIPL this is already on – don't touch**

13.2. *mem_pll_initf* : PLL Initfile for MBAs

a) *cen_mem_pll_initf.C*

- MBA PLL setup – get tune bits, frequency from attributes
 - Note that Hostboot doesn't support twiddling bits, looks up which “bucket” (ring) to use from attributes set during *mss_freq*
 - Data is stored as a ring image that is frequency specific
 - 6 different frequencies – 6 different ring images
- Must run at system frequency

13.3. *mem_pll_setup* : Setup PLL for MBAs

a) *cen_mem_pll_setup.C*

- MBA PLL setup
 - Moved PLL out of bypass(just DDR)
- Performs PLL checking

13.4. *mem_startclocks* : Start clocks on MBAs

a) *cen_mem_startclocks.C*

- Drop fences and tholds on MBAs

13.5. *host_enable_vddr* : Enable the VDDR3 Voltage Rail

a) Bring power to dram. Must be done here to meet JEDEC spec compliance

- **On MPIPL this is already on – don't touch**
- b) For OpenPower systems the host directly powers on the VDDR
 - This can be many mechanisms, but primary one is via a I2C attached 955x GPIO control
 - Optionally at this point Hostboot can tweak the VRMs to compensate for number of plugged dimms

13.6. *mss_scominit* : Perform scom inits to MC and PHY

a) *mss_scominit.C*

- Called per Centaur, pass in associated MBAs

- HW units included are MBS, MBX, MBAs, and PHYs
- Initfiles: cen_ddrphy.initfile, cen_mba.initfile, cen_mbx.initfile and cen_mbs.initfile
- HW calls multiple initfiles per target
- FAPI provides scom initfile handling
- Uses attributes from previous step

b) *proc_throttle_sync.C*

- Must be issued on all P8s, can only be issued after ALL centaurs on given p8 have scominit complete (can also loop at the end)
- Triggers sync command from MCS to actually load the throttle values into the centaurs

13.7. *mss_ddr_phy_reset* : Soft reset of DDR PHY macros

a) *mss_ddr_phy_reset.C* : Soft Reset the DDR phy macros

- DDR PLLs
 - Already configured DDR PLL in scaninit
 - Take PLL out of reset (controlled by a GP3 bit)
- Sends Soft DDR Phy reset
- Kick off internal ZQ Cal
- Perform any config that wasn't scanned in (TBD)
 - Nothing known here

13.8. *mss_draminit* : Dram initialize

a) *mss_draminit.C* : Init the DDR logic

- De-assert dram reset
- De-assert bit (Scom) that forces mem clock low – dram clocks start
- Raise CKE
- Load RCD Control Words
- Load MRS – for each dimm pair/ports/rank
 - ODT Values
 - MR2, MR3, MR1, MR0

13.9. *mss_draminit_training* : Dram training

a) *mss_draminit_training.C*

- Prior to running this procedure will apply known DQ bad bits to prevent them from participating in training. This information is extracted from the bad DQ attribute and applied to Hardware
 - Marks the calibration fail array
- External ZQ Calibration
- Execute initial dram calibration (7 step – handled by HW)
- This procedure will update the bad DQ attribute for each dimm based on its findings

13.10. *mss_draminit_trainadv* : Advanced dram training

a) *mss_draminit_training_advanced.C*

- Prior to running this procedure will apply known DQ bad bits to prevent them from participating in training. This information is extracted from the bad DQ attribute and applied to Hardware

- Marks the MCBist mask
 - This step will contain any algorithms to improve eye
 - Also will contain some characterization (mfg only) tests
 - All optional tests must be their own functions and C files to optimize load and execution
 - This procedure will update the bad DQ attribute for each dimm based on its findings
- 13.11. `mss_draminit_mc` : Hand off control to MC
- a) `mss_draminit_mc.C`
- Set IML complete bit in centaur
 - Start main refresh engine
 - Refresh, periodic calibration, power controls
 - Unmask memory FIRs
 - Turn on ECC checking on memory accesses
- b) Note at this point memory FIRs can be monitored by PRD
- 13.12. `mss_dimm_power_test` : Hand off control to MC
- a) `mss_dimm_power_test.C`
- Pass in all dimms on a given power domain
 - Will be a no-op for CDIMMs (FW can always call)
 - For IS dimms measure the power
14. **Step 14 Hostboot – DRAM Initialization**
-
- 14.1. `host_startprd_dram` : Load PRD for DRAM domain
- a) Expand Host PRD to include DRAM (as well as powerbus and DMI)
- b) Must be minimal environment (running out of L3)
- 14.2. `mss_extent_setup`
- a) `mss_extent_setup.C`
- Pushes memory extent configuration into the centaurs via Inband
 - Addresses are pulled from attributes, set previously by `mss_eff_config`
 - Centaurs always start at address 0, address map controlled by `proc_setup_bars` below
- 14.3. `mss_memdiag` : Mainstore Pattern Testing
- a) The following step documents the generalities of this step
- In FW PRD will control mem diags via interrupts. It doesn't use `mss_memdiags.C` directly but the HWP subroutines
 - In conus it will execute `mss_memdiags.C` directly
- b) `mss_memdiags.C`
- Prior to running this procedure will apply known DQ bad bits to prevent them from participating in training. This information is extracted from the bad DQ attribute and applied to Hardware
 - Will use traditional scrub engine
 - Still supports superfast
 - Modes:
 - Minimal: Write-only with 0's
 - Medium: Write-followed by Read, 4 patterns, last of 0's
 - Max: Write-followed by Read, 9 patterns, last of 0's
 - Final “init” done with scrub engine to ensure good ECC

- Note that this does not need to be zeros, but it is typically zeroed (PHYP will zero again)
 - Run on the host
 - This procedure will update the bad DQ attribute for each dimm based on its findings
 - At the end of this procedure sets FIR masks correctly for runtime analysis
- c) All subsequent repairs are considered runtime issues
- 14.4. `mss_thermal_init` : Initialize the thermal sensor
- a) `mss_thermal_init.C`
- Called on Centaur target
 - Setup and configure I2C thermal sensor on dimms
 - Configure and start centaur thermal cache
 - Configure and start the OCC cache
 - Disable safe mode throttles
 - Will cause memory to go to runtime emergency throttles
 - When OCC starts polling OCC cache will revert to runtime settings
- b) `proc_throttle_sync.C`
- Must be issued on all P8s, can only be issued after ALL centaurs on given p8 have thermal init complete (can also loop at the end of all centaurs)
 - Triggers sync command from MCS to actually load the throttle values into the centaurs
- 14.5. `proc_pcie_config` : Configure the PHBs
- a) `proc_pcie_config.C`
- Called on all chips, target is per PHB
 - Procedural based – will call initfile if need be
 - Covers PCIe Phase 2 Inits 18-30
 - Setup config regs
 - Command and Data credits
 - Clear FIRs (if needed)
 - Unmask PCIe FIRs
- 14.6. `mss_power_cleanup` : Clean up any MCS/Centaurs
- a) `mss_power_cleanup.C`
- Called on all present Centaurs and MBAs
 - Cleans up and powers down unused centaurs/mcs/DMI
 - **After this step no more HW reconfig loops (back to step 10) are allowed – this step and future ones will not perform the reconfig loop check**
 - Hostboot will start to flow out to memory in the next step
 - Any memory errors after this point are considered “runtime errors”
 - All errors from this point on have to be a no deconfig and guard OR terminate the IPL (and let the SP do the reconfig)
 - If user attempts to do a deconfig outside the loop – then attempt to fail
- 14.7. `proc_setupBars` : Setup Memory BARs
- a) `mss_setupBars.C`

- Based on system memory map
 - Each MCS has its mirroring and non mirrored BARs
 - Set the correct checkerboard configs. Note that chip flushes to checkerboard
 - need to disable memory bar on slave otherwise base flush values will ack all memory accesses

b) *proc_setup_bars.C*

- Sets up Powerbus/MCD, L3 BARs on running core
 - Other cores are setup via winkle images
- Setup dSMP and PCIe Bars
 - Setup PCIe outbound BARS (doing stores/loads from host core)
 - Addresses that PCIe responds to on powerbus (PCI init 1-7)
 - Informing PCIe of the memory map (inbound)
 - PCI Init 8-15
- Set up Powerbus Epsilon settings
 - Code is still running out of L3 cache
 - Use this procedure to setup runtime epsilon values
 - Must be done before memory is viable

14.8. *proc_exit_cache_contained* : Allow execution from memory

a) *proc_exit_cache_contained.C*

- Allow execution to flow out to memory
- Must clear CoreFabricProtect bit
- Must clear NestFabricProtect bit
- Must set all centaur's secureboot bits to prevent memory D/A

b) Data rolls out to memory

14.9. *host_mpipl_service* : Perform MPIPL tasks

- a) This is a no-op for warm/cold IPLs. See description in Error: Reference source not found, Section Error: Reference source not found for full details

15. Step 15 Hostboot – Build Winkle Images

15.1. *host_build_winkle* : Build runtime winkle images

a) Pull Reference Image from PNOR

- Run through secure boot algorithm

b) *p8_slw_build.C*

- Called for each processor chip.
- Parameter: Pointer to Reference image.
- Parameter: Pointer to Output winkle image.
 - This is any Hostboot specified mainstore location (does not have to be attached to the processor being winkled).
 - When PHYP is loaded, the winkle images will be trampled, PHYP will call p8_slw_build to recreate them in a PHYP specified location in mainstore (each image will probably be placed in mainstore local to its associated processor for performance).
- Customize image with data for each core
 - Scan rings – Time, GPTR, Repair

- Tweak to make runtime acceptable – expect to be only scom registers
 - Write image to output pointer parameter
- c) *p8_pore_gen_cpu_reg()*
- API that updates a winkle image with various chip state registers (MSR, HRMOR, LPCR)
 - The chip registers are set to these values on winkle exit
 - This will only be called by Hostboot. Cronus will not use it. Hence separate from *proc_slw_build*.
- 15.2. *proc_set_pore_bar* : Tell SWL Eng where winkle image is
- a) Cronus will load the images via procpupmem
- b) *p8_set_pore_bar.C*
- Called for each processor chip.
 - Parameter: Physical address of SLW image (not a pointer address, it cannot be dereferenced)
 - Parameters: PBA BAR number, SLW image size, SLW image location (default: mem; others: L3, SRAM)
 - *p8_pba_bar_config.C (called as subroutine)*
 - Set BAR address
- 15.3. *host_poreslw_init -init* : Initialize the PORE-SLW engine and related functions to allow assisted idles to operate
- a) *p8_poreslw_init.C chip_target, ENUM:PM_INIT*
- Called for each processor chip
 - Parameters: none other than target
 - Set and then clear PMC reset to clear any latent idle transitions that may be queued (in the case this flow is executed as part of a warm IPL or MPIPL).
 - **Note** this clears ALL PMC settings. However, previously issued SPIVID operations will remain intact.
 - *p8_pmc_deconfig_setup.C (called as subroutine)*
 - Reads the EX chiplet enable (GP0(0)) and sets the PMC_CORE_DECONFIGURATION_REG for any non-configured EX chiplets
 - This is necessary to enable the PMC to accept idle entry/exit requests from on the validly enabled PCBS. This setting is also necessary for proper operation of the Pstate mechanism.
 - Make sure the PMC Idle Sequencer is enabled
 - Clears PMC_MODE_REG(14) - HALT_IDLE_STATE_MASTER_FSM
 - Enables OHA to accept idle operations
 - Clear OHA_MODE_REG(6) - OHA Idle State Override Enable
 - Activates the PCBS-PM macro
 - Clear PMGP0(0) (PM_DISABLE) as this is necessary for both idle and Pstate operation. Hardware default is to come up inactive.
 - Clear the OCC Special Wake-up which is initialized ON by hardware default to keep idle operations from occurring until the SLW image is installed.

16. Step 16 Hostboot – Core Activate

16.1. *host_activate_master* : Activate master core

a) *proc_prep_master_winkle.C*

- This procedure checks to make sure that the SBE program to exit winkle is already running from PIBMEM (left running from initial SBE steps). For both real SBE and vSBE it triggers deadman

loop

- ***proc_sbe_trigger_winkle.S*** (loaded into *PIBMEM* prior or via *vSBE*)
 - Monitor master winkle and trigger winkle exit when it detects winkle enter. (Note that this is running in parallel with Hostboot, started at end of istep 5)
 - Starts the deadman timer loop
 - Sets indication that SBE program is ready
 - Waits for winkle entered (completely entered)
 - SBE triggers winkle exit via IPI to all threads on a core. Uses istep control bits to only start specific threads if in AVP mode
 - If Hostboot does not stop deadman timer in 10 seconds checkstop system

b) *proc_trigger_winkle* – Hostboot path (Hostboot running)

- Hostboot function, not a HW Procedure
- ***p8_block_wakeup_intr.C-set***
 - This will prevent all interrupts/wake up sources to the core, thus allowing the next step (winkle) to work
- Issue system call to cause all threads to enter winkle. Core will then enter winkle state
 - Clear LPCR (cover not entering due to external interrupts)
 - issue winkle instruction

c) *proc_force_winkle.C* – SP/Cronus path (Hostboot not running)

- Since Hostboot is not running this procedure will force the master core to winkle
 - Master core is left in initialed, maintenance mode state after proc SBE

d) *proc_trigger_winkle.C* – Cronus path only

- This procedure is a NO-OP when the real SBE is executing. It is hook to allow the virtual SBE to trigger the winkle exit – ie resume execution of ***proc_sbe_trigger_winkle.S*** in vSBE mode (see above)
 - Note that this is NOT supported in SP based IPL, Cronus only

e) *proc_stop_deadman_timer.C*

- Hostboot runs when active, otherwise Cronus will have to execute
 - Stops the deadman timer
 - Stops the SBE program

16.2. *host_activate_slave_cores* : Activate slave cores

a) Hostboot active:

- Setup stack space for all slave core threads –
- Wake up all threads on all cores
 - Cores are sitting in a winkled state (flush that way)
 - Issue IPI to all slave cores to force winkle exit. Will start executing at SREST vector (0x100). Bring them into Hostboot collective

b) Hostboot not running:

- Cores come alive and into maintenance mode (LPCR not set)
- ***proc_trigger_winkle.C***
 - Called on a core target
 - SP/Cronus issue IPIs to all cores/threads in system except for those on master core

16.3. `mss_scrub` : Start background scrub

a) `mss_scrub.C`

- Note that this is not executed directly by Hostboot (instead triggered by PRD), Cronus will execute HWP directly
- Start background scrubbing. First pass will trigger as fast as possible, then switch to 12h scrub cycle
- Currently Hostboot will not wait (block) before flowing out to memory
- The completion of the scrub commands must be handled by SP or (PHYP/Sapphire) PRD
- HostPRD will not be called after this point (not called for this step)

16.4. `host_ipl_complete` : Notify SP drawer ipl complete

a) Stop hostPRD (in anticipation that SP will take over PRD responsibilities)

- For OpenPower systems PRD will resume once OPAL is alive

b) `cen_switch_rec_attn.C`

- Switch recoverable/special attentions from host back to FSP

c) `proc_switch_rec_attn.C`

- Switch recoverable/special attentions from host back to FSP

d) `proc_switch_cfsim.C`

- Switch the FSI bus control back to the FSP
- Needed for Secureboot

17. Step 17 FSP – Init PSI

17.1. No-op on OpenPower systems

18. Step 18 Establish System SMP and TOD clocks

18.13. `proc_tod_setup`

a) Enable the PCIe OscLite macro – NOT called on MPIPL

- `proc_enable_osclite.C`
 - Turn off the 'power-pon-reset' to osclite macro
 - Setup oscillator mode based on istep 0 setup
 - Check that osclite matches expected output (if not returns an error for FW to trigger reconfig loop)

b) FW owns algorithm of TOD topology, HWP pushes values into HW

c) `proc_tod_setup.C`

- FW passes in a topology tree, which TOD oscillator to use, and primary/secondary topology
- HWP determines delay values from attributes (MRW)
- HWP programs HW
- HWP outputs register values needed for PHYP and PRD analysis

18.14. `proc_tod_init`

a) Performed to init the TOD network. Done during the FW IPL due to AVPs, note that it will be done again by PHYP when they start

- `proc_tod_init.C`
- Setup EX chiplet TOD

- b) This is the last istep that FSP/HWSV will perform a HW reconfig loop

19. Step 19 SP – Prepare for Host

- 19.1. No-op for OpenPower systems

20. Step 20 Hostboot – Load Payload

- 20.1. No-op for OpenPower systems

21. Step 21 Hostboot – Start Payload

21.1. `host_runtime_setup`

- a) Note that this step is only issued to master HB instance
- b) Take down any/all TCE setup
- c) For PHYP based systems
 - Loop through attributes and write them to predefined memory area inside of the HDAT structures
 - Note: HB master issues IPC to HB slaves for them to update their sections
 - Append the TPM log to HDAT structures
 - Note: HB master issues IPC to HB slaves for them to update their sections
- d) For OPAL based systems
 - Loop through attributes and build a device tree for OPAL
- e) In AVP or OPAL mode Hostboot will load the OCC and start it here. If the load/start fails then HB will send a errorlog to the SP and the SP will terminate the IPL
 - OCC must monitor for the broadcast scom read (OR) of EX scratch register 7 for the removal of the payload started signature before using the FSI2Host mailbox for ATTN traffic. Note that OCCs on non master chips will never have to wait (as Hostboot only uses the FSI2Host mailbox on the master chip)

21.2. `host_verify_hdat`

- a) Only issued to master HB instance
 - If needed IPC to slaves to perform their tasks
- b) Secureboot verification of PHYP/AVP image load

21.3. `host_start_payload`

- a) Prior to starting shutdown sequence Hostboot must write hostboot (ASCII) to scratch register 7 on the master core. All other cores on the master chip must be written to same value or 0s. This value will be polled by the FSP in the next step to ensure that hostboot has truly quiesced
- b) Hostboot enters shutdown sequence
 - Quiesce mailbox and all DMAs
 - Flush data to PNOR
 - Disable interrupts
 - Send sync message to FSP (or respond to istep)
 - Enter Kernel
 - Prepare to jump to payload – at this point hostboot must not TI
 - Clear scratch register 7 on master core
- c) Payload is started by

- switching HRMOR to desired address and jumping to entry point
 - Note that master thread must be the last one to jump
 - payload cannot start until all threads have been transitioned
 - For multi-node systems the HB master does the following:
 - Issue slave node shutdown request via IPC
 - HB master polls the “Hostboot done scratch reg” for all slave nodes to enter payload
 - HB Master issues own shutdown
- d) No Hostboot code is reused, only mechanism is data passed in HDAT areas. Hostboot runtime is a part of PHYP

5 Hostboot Runtime Services (HBRT)

The following are not IPL time procedures, but functions called by PHYP or OPAL on Hostboot to perform various tasks. The numbering has been kept common (for convention), but they are not guaranteed to run in this order

State at this point

- PHYP/OPAL running
- Memory is initialized
- SMP alive
- All cores have gone through winkle and are running

22. Enable Winkle

This step is controlled and issued by PHYP/OPAL when they are ready to build the winkle image. It is required that the winkle image be present in memory prior to loading and starting the OCC.

22.1. `host_build_winkle` : Build runtime winkle images

a) Pull Reference Image from SP or PNOR for OpenPower systems

- Run through secure boot algorithm

b) `p8_slw_build.C`

- Called for each processor chip.
- Parameter: Pointer to Reference image.
- Parameter: Pointer to Output winkle image.
 - This is any Hostboot specified mainstore location (does not have to be attached to the processor being winkled).
 - When PHYP is loaded, the winkle images will be trampled, PHYP will call `p8_slw_build` to recreate them in a PHYP specified location in mainstore (each image will probably be placed in mainstore local to its associated processor for performance). OPAL leaves the original images in place from Hostboot
- Customize image with data for each core
 - Scan rings – Time, GPTR, Repair
 - Tweak to make runtime acceptable – expect to be only scom registers
- Write image to output pointer parameter

22.2. `proc_set_pore_bar` : Tell SWL Eng where winkle image is

a) *p8_set_pore_bar.C*

- Called for each processor chip.
- Parameter: Physical address of SLW image (not a pointer address, it cannot be dereferenced)
- Parameters: PBA BAR number, SLW image size, SLW image location (default: mem; others: L3, SRAM)
- *p8_pba_bar_config.C (called as subroutine)*
 - Set BAR address

22.3. *p8_poreslw_init -init* : Initialize the PORE-SLW engine and related functions to allow assisted idles to operate

a) *p8_poreslw_init.C chip_target, ENUM:PM_INIT*

- Called for each processor chip
- Parameters: none other than target
- Set and then clear PMC reset to clear any latent idle transitions that may be queued (in the case this flow is executed as part of a warm IPL or MPIPL).
- **Note** this clears ALL PMC settings. However, previously issued SPIVID operations will remain intact.
- *p8_pmc_deconfig_setup.C (called as subroutine)*
 - Reads the EX chiplet enable (GP0(0)) and sets the PMC_CORE_DECONFIGURATION_REG for any non-configured EX chiplets
 - This is necessary to enable the PMC to accept idle entry/exit requests from on the validly enabled PCBS. This setting is also necessary for proper operation of the Pstate mechanism.
- Make sure the PMC Idle Sequencer is enabled
 - Clears PMC_MODE_REG(14) - HALT_IDLE_STATE_MASTER_FSM
- Enables OHA to accept idle operations
 - Clear OHA_MODE_REG(6) - OHA Idle State Override Enable
- Activates the PCBS-PM macro
 - Clear PMGP0(0) (PM_DISABLE) as this is necessary for both idle and Pstate operation. Hardware default is to come up inactive.
- **Clear the OCC Special Wake-up which is initialized ON by hardware default to keep idle operations from occurring until the SLW image is installed**

p8_pore_gen_cpu_reg() will be called by PHYP/OPAL prior to winking any core

- API that updates a winkle image with various chip state registers (MSR, HRMOR, LPCR)
 - The chip registers are set to these values on winkle exit
 - **This will only be called directly by PHYP/OPAL at their discretion. Hence separate from *proc_slw_build*.**

23. Reset and Initialize OCC

Reordered to have OCC/PM hardware initialized first and then have the OCC image loaded in memory to allow for separation of where the functions run. Due to Trusted Boot, the hardware initialization procedures need to run under Hostboot

This step will run each of the substeps to each chip within a physical node (an OCC boundary) before proceeding to the next step. This is done as a regular process in looking to the the “start_occ” step whereby the OCCs will start in reasonable time proximity (one followed by the next via singular XSCOM to each chip) to minimize OCC startup timeouts.

23.1. Setup OCC bars : Establish legal addressing

a) *p8_pba_bar_config.C chiptarget, address*

- Address dictated by PHYP
- Called once for each of 4 BARs
- Place image in EM Nodal Region at offset 0

23.2. power_management_reset : Reset Power Management (includes clearing any latent errors that may be pending; done for the case of OCC reset)

a)

For the Murano DCM, one chip controls the voltage (via SPIVID) used by both chips and an on-module interface exists between the PMCs of the two chips to 1) coordinate Pstate operations that lead to voltage changes; and 2) communication between the OCCs on each chip. Therefore, when performing a reset, operations must happen to each chip in a coordinated manner so that the logic does not land in unrecoverable states brought about by asynchronous reset operation. The PMC implements a set of “halt” operations that allow for on-going operations to complete naturally unless prevented by flagged errors. This was done to reduce the reset state space.

In order to perform operations to these two chips, the relevant procedures need to have chip_targets for both chips available. The caller must determine what the targets are for the chips that are physically in the DCM. The “tie” is the PowerBus NodeID as processors attached by X-Buses are in the same node.

For Venice (where DCMs are not supported), a NULL secondary target is to be passed.

b) *p8_pm_prep_for_reset.C *primary_chiptarget, *secondary_chiptarget*

- For primary_target and then secondary target (if not NULL),
- *p_occ_control.C *chiptarget, ENUM:OCC_STOP*
 - OCC PPC405 put into reset
 - Chiplets forced (in PCBS) to Psafe value due to heartbeat loss
 - If chiplet is running above this, will force a clip (independent of Global Actual)
 - If chiplet is running below this, will cause a rise in frequency to match the current Global Actual
 - This also cuts of an new requests (PHYP if so enabled) from influencing frequency
 - PMC moves to Vsafe value due to heartbeat loss
 - This will cause the external rails (Global Actual) to a value consistent with Psafe
- *p8_cpu_special_wakeup.C *chiptarget, ENUM:OCC_SPECIAL_WAKEUP*
 - (Not used by PHYP – custom procedure used)
 - For each chiplet, put into Special Wake-up via the OCC special wake-up bit.
 - Doesn't collide with FSP/PHYP bit.
 - Takes the PORE-SW, OCI and PBA out of the equation
 - Take PCBS PFET controller out of the equation
 - Poll for completion.
 - If timeout, indicate that restart of OCC is to not occur via fapi::ReturnCode
 - RC_PROCPM_SPC_WAKEUP_TIMEOUT
 - PRD effect: Mark chiplet for garding

- c) Note: PORE-SLW detected errors will produce malfunctions alerts to PHYP whereby the set of events defined in *p8_poreslw_recovery.C* occur to deal with getting the idle handling complex recovered for use.
- *p8_pmc_force_vsafe.C *chiptarget*,
 - Forces the Vsafe value into the voltage controller. This will cease processing of any new PState requests and make the Global Actual take on the known value unless the voltage controller is in error.
 - Note: setting this in the “master” chip (the one with SPIVID attached and which may be either the primary or the secondary target) will cause the Pstate process to forcefully move. Setting this into the “slave” chip simply forces its contribution to the Global Pstate (as a request) to the Pvsafe value. Given that they are configured to the same Pvsafe value (which is done in the -init phase), the necessary end result occurs (really driven by the forcing in the “master” chip) no matter which order the chips are targeted.
 - Value is contained in PMC_PARAMETER_REG1.Pvsafe
 - **Implementation: force PMC OCC Heartbeat to 0 to make the hardware react**
 - Check PMC state/status that Vsafe (Pstate) as been achieved.
 - If timeout, indicate that restart of OCC is to not occur (means are TBD)
 - Else (the slave chip)
 - Exit with success as the Master is in control of the voltage.
- *p8_pcbs_init.C *chiptarget, ENUM:PM_RESET*
 - For **all** configured EX chiplets
 - Force safe mode (uses Psafe Pstate setting) in EX chiplet (multicastable) (for safety) (with removal of MIN functions (design change))
 - Check PCBS-PMPM state/status that Psafe (Pstate) as been achieved and that FSM are in a stable state.
 - If timeout, indicate that restart of OCC is to not occur (means are TBD) Else
 - Write DPLL Frequ Control Reg with the frequency (not PState!) that represents the minimum of Psafe or Global Actual as seen in the PCBS.
 - Chiplets forced (via PCBS) into DPLL Override / PState disabled mode (whatever the above PState turned out to be)
 - The frequency value determined by Psafe/Vsafe is locked from changes by the PState mechanism. This allows the PCBS to be re-initialized.
 - Force OCC SPR Mode in EX chiplet (multicastable)
 - Force *global_en PState to off to cease interrupts to PMC (Global, AutoOverride, Idle) in EX chiplet (multicastable)
 - Reset Pmin and Pmax to -128 (0x80) and 127 (0x7F) in EX chiplet (multicastable)
 - Chiplets resonant clocking (via PCBS) disabled
 - This allows the PCBS to be re-initialized blindly
 - **It is debatable whether this is necessary as resonant clock can continue to operate while the OCC reset is being performed and while the chiplet is running at Psafe).**
 - OCC Heartbeat disable
 - Will be enabled by OCCFW (not FAPI)
 - Chiplet IVRMs put into bypass mode and then disabled
 - This allows procedures to reload the Local PState and VDS Tables

- Undervolting values reset: Kuv to 0; Puv_min and Puv_max each set to -128 each (to disable)
- Local Pstate Frequency Target mechanism disabled
 - OCCFW will enable it
- Issue reset to PCBS-PM (dial endp_reset_pm_only)
 - DPLL Frequ Control Reg is NOT reset
 - Scan only value is 1GHz
- ***p8_pmc_init.C *primary_chiptarget, *secondary_chiptarget, ENUM:PM_RESET***

This is meant to primarily reset the PMC for restarting the OCC for firmware purposes (reset due to firmware error and firmware update). As it will reset the PMC hardware to it IPL state, it has also a high probability of clearing hardware detected errors. However, hardware errors can be very large state space so the PMC design philosophy (due to the primary purpose above) is to gracefully halt and check for a good halt condition.

Actions to a *secondary_chiptarget* in the following are only done if the **secondary_chiptarget* is not NULL. Consistency checking will be performed using the value is the attribute ATTR_DCM_INSTALLED. If this attribute is set to “true”, then **secondary_chiptarget* is expected to be non-NULL.

This procedure can determine which of the passed targets is the DCM “master” by checking if the SPIVID is enable for that target. In this way, it can determine the mapping of the primary/secondary to master/slave for the order noted below.

Prerequisite: the OCC (PPC 405) is in the halted state when this procedure executes

- Issue halt to Pstate Master FSM on *master_chiptarget*
- Issue halt to Pstate Master FSM on *slave_chiptarget*
- Issue halt to interchip FSM on *master_chiptarget*
- Poll for interchip interface to stop on *master_chiptarget*
 - If poll not complete, flag “reset_suspicious” and save the poll point; continue
- Issue halt to interchip FSM on *slave_chiptarget*
- Poll for interchip interface to stop on *slave_chiptarget*
 - If poll not complete, flag “reset_suspicious” and save the poll point; continue
- If voltage changes are enable, issue halt to SPIVID controller on FSM on *master_chiptarget*
- Poll for SPIVID FSM to halt on *master_chiptarget*
 - If poll not complete, flag “reset_suspicious” and save the poll point; continue
- Poll for Pstate Master FSM being stopped on *master_chiptarget*
 - If poll not complete, flag “reset_suspicious” and save the poll point; continue
- Poll for Pstate Master FSM being stopped on *slave_chiptarget*
 - If poll not complete, flag “reset_suspicious” and save the poll point; continue
- Poll for O2P bridge being complete on *master_chiptarget*
 - As the OCC (PPC405) has been halt prior to entry of this procedure, this poll will be immediate if there are no PIB related errors present
- Poll for O2P bridge being complete on *slave_chiptarget*
 - As the OCC (PPC405) has been halt prior to entry of this procedure, this poll will be immediate if there are no PIB related errors present
- Poll for O2S bridge being complete on *master_chiptarget*

As the OCC (PPC405) has been halt prior to entry of this procedure, this poll will be immediate if there are no SPIVID related errors present

- Poll for Idle FSM being quiesced (timeout: 500ms to cover the case of having all 4 types of Deep Idle transitions in flight)

Note: Previously issued special wake-ups could have triggered PORE activity through the Idle FSM (and the related pending queues).

- if poll timeout, mark the error point
- Note on Idle/PORE-SLW state (prior to reset)

Given that special wake-up occurred before this point, any errors that resulted from that special wake-up (eg PORE-SLW fatal or timeout indicated in PMC LFIR) will have fired a malfunction alert to PHYP whereby the execution of **p8_poreslw_recovery.C** will have taken place.

- Issue reset to the PMC

Note: this action will wipe out the Idle Pending queue so that requests for idle transitions (entry and exit) will be lost which means that PHYP notification needs to happen.

- Write **PMC_MODE_REG.pmc_reset_all_voltage_registers = 1**.
 - Clearing LFIRs will have been done by PRD
 - Note: this will remove CONFIG settings
 - This puts the PMC into firmware mode which halts any future Global Actual operations
- For primary_target and then secondary target (if not NULL),
 - **p8_poreslw_init.C *chiptarget, ENUM:PM_RESET**
 - Issue reset to PORE Sleep/Winkle engine
 - With Special Wake-up in place, this engine is not being used.
 - **p8_poregpe_init.C *chiptarget, ENUM:PM_RESET**
 - Issue reset to PORE General Purpose Engine
 - With PPC405 held in reset, this engine is not being used.
 - **p8_oha_init.C *chiptarget, ENUM:PM_RESET**
 - TODO: Need to check on AISS reset
 - **p8_pba_init.C *chiptarget, ENUM:PM_RESET**
 - Issue any operations to return FSMs to idle
 - **p8_occ_sram_init.C *chiptarget, ENUM:PM_RESET**
 - Issue any operations to return FSMs to idle
 - **p8_ocb_init.C *chiptarget, ENUM:PM_RESET**
 - Disable all OCB indirect channels and return them to their power-on state
 - Channels 0-2:
 - OCBCSRn: 0 to WAND to clear errors, 0x04000000 to WOR to set Circular mode
 - OCBESRn: 0 to clear errors
 - OPPCINJ: 0 to clear injection modes
 - OCBSLCSn: 0 to disable the pull queue and all settings
 - OCBSLBRn: 0 the region and pull start address
 - OCBSHCSn: 0 to disable the push queue and all settings
 - OCBSHBRn: 0 the region and push start address

- OCBSESn: 0 to clear errors
- OCBICRn: 0 to reset “allow_unsecure_pib_masters”
- OCBLWCRn: 0 to disable the linear window and clear the BAR/Mask values
- OCBLWSBRn: 0 the linear window region and linear window base
- Channel 3:
 - OCBCSRn: 0x00000000 to WAND to clear errors, set Linear mode
 - OCBESRn: 0 to clear errors
 - OCBSESn: 0 to clear errors
 - OCBICRn: 0 to reset “allow_unsecure_pib_masters”
- Remove error injection modes
 - OPPCINJ: 0 to clear injection modes
- ***p8_pss_init.C *primary_chiptarget, *secondary_chiptarget, ENUM:PM_RESET***
 This procedure determines if either of the targets passed has a SPIPSS attached by looking for the SPIPSS being enabled for either target. In this way, it can determine the mapping of the primary/secondary to the one that will perform the operations below. If neither chip has a SPIPSS attached, then the procedure simply exits with success.
- See that any outstanding operations have finished in ADC engine
 - Read, modify write
 SPIPSS_ADC_CTRL_REG1.enable_restart_ADC_sampling_after_last_read = 0;
 disable automatic resample while retaining CONFIG settings
 - Poll for 1ms on SPIPSS_ADC_STATUS_REG.hwctrl_ongoing being 0
- See that any outstanding operations have finished in P2S engine
 - Poll for 1ms on SPIPSS_P2S_STATUS_REG.p2s_ongoing being 0
- Write to SPIPSS_ADC_RESET_REGISTER and SPIPSS_PSS_RESET_REGISTER to perform the reset
 - **Note: this will remove CONFIG settings**

23.3. power_management_Init : Initialize Power Management

a) ***p8_pm_init.C *primary_chiptarget, *secondary_chiptarget, ENUM:PM_CONFIG***

Calls the underlying unit procedures (***p8_<unit>.init.C with ENUM:PM_CONFIG***) to read any platform attributes and create necessary feature attributes that will be used by the ***p8_<unit>.init.C with ENUM_PM_INIT***

For a DCM, the order of operations between the master/slave chips (physical) vs the primary/secondary (logic) is not important in the phase.

- For primary_target and then secondary target (if not NULL),
 - ***p8_pcbcs_init.C *chiptarget, ENUM:PM_CONFIG***
 - PFET Sequencing Delays
 - **convert_pfet_delays()** - Convert the following delays from platform attributes (binary in nanoseconds) to PFET delay value feature attributes. The conversion uses **ATTR_PROC_NEST_FREQUENCY**.
 - Input platform attributes
 - **ATTR_PM_PFET_POWERUP_CORE_DELAY0**
 - **ATTR_PM_PFET_POWERUP_CORE_DELAY1**
 - **ATTR_PM_PFET_POWERUP_ECO_DELAY0**

- ATTR_PM_PFET_POWERUP_ECO_DELAY1
- ATTR_PM_PFET_POWERDOWN_CORE_DELAY0
- ATTR_PM_PFET_POWERDOWN_CORE_DELAY1
- ATTR_PM_PFET_POWERDOWN_ECO_DELAY0
- ATTR_PM_PFET_POWERDOWN_ECO_DELAY1
- Output feature attributes
 - ATTR_PM_PFET_POWERUP_CORE_DELAY0_VALUE
 - ATTR_PM_PFET_POWERUP_CORE_DELAY1_VALUE
 - ATTR_PM_PFET_POWERUP_CORE_SEQUENCE_DELAY_SELECT
 - ATTR_PM_PFET_POWERUP_ECO_DELAY0_VALUE
 - ATTR_PM_PFET_POWERUP_ECO_DELAY1_VALUE
 - ATTR_PM_PFET_POWERUP_ECO_SEQUENCE_DELAY_SELECT
 - ATTR_PM_PFET_POWERDOWN_CORE_DELAY0_VALUE
 - ATTR_PM_PFET_POWERDOWN_CORE_DELAY1_VALUE
 - ATTR_PM_PFET_POWERDOWN_CORE_SEQUENCE_DELAY_SELECT
 - ATTR_PM_PFET_POWERDOWN_ECO_DELAY0_VALUE
 - ATTR_PM_PFET_POWERDOWN_ECO_DELAY1_VALUE
 - ATTR_PM_PFET_POWERDOWN_ECO_SEQUENCE_DELAY_SELECT
- *p8_poreslw_init.C *chiptarget, ENUM:PM_CONFIG*
 - None defined
- *p8_poregpe_init.C *chiptarget, ENUM:PM_CONFIG*
 - None defined
- *p8_oha_init.C *chiptarget, ENUM:PM_CONFIG*
 - **convert_ppt_time()** - Convert Power Proxy Trace Time to Power Proxy Trace Time Select and Match feature attributes
 - With ATTR_PM_POWER_PROXY_TRACE_TIMER (binary in nanoseconds) to produce ATTR_PM_PPT_TIMER_MATCH_VALUE and ATTR_PM_PPT_TIMER_TICK
 - 0=0.25us , 1=0.5us, 2=1us, and 3=2us
- *p8_pmc_init.C *chiptarget, ENUM:PM_CONFIG*
 - **pmc_create_spivid_settings()** - Translate
 - Input platform attributes
 - ATTR_PROC_NEST_FREQUENCY
 - ATTR_PM_SPIVID_FREQUENCY
 - ATTR_PM_SPIVID_FRAME_SIZE
 - ATTR_PM_SPIVID_IN_DELAY_FRAME1
 - ATTR_PM_SPIVID_IN_DELAY_FRAME2
 - ATTR_PM_SPIVID_INTER_FRAME_DELAY
 - ATTR_PM_SPIVID_INTER_FRAME_DELAY_WRITE_STATUS
 - ATTR_PM_SPIVID_INTER_RETRY_DELAY
 - Output feature attributes

- ATTR_PM_SPIVID_CLOCK_DIVIDER
 - ATTR_PM_SPIVID_OUT_COUNT
 - ATTR_PM_SPIVID_IN_COUNT
 - ATTR_PM_SPIVID_INTERFRAME_DELAY_SETTING
 - ATTR_PM_SPIVID_INTERFRAME_DELAY_WRITE_STATUS_VALUE
 - ATTR_PM_SPIVID_INTER_RETRY_DELAY_VALUE
 - *p8_pba_init.C *chiptarget, ENUM:PM_CONFIG*
 - None defined
 - *p8_occ_sram_init.C *chiptarget, ENUM:PM_CONFIG*
 - None defined
 - *p8_och_init.C *chiptarget, ENUM:PM_CONFIG*
 - None defined
 - *p8_pss_init.C *chiptarget, ENUM:PM_CONFIG*
 - create_spipss_settings() - Translate
 - Input platform attributes
 - ATTR_PROC_NEST_FREQUENCY
 - ATTR_PM_SPIPSS_FREQUENCY
 - ATTR_PM_SPIPSS_FRAME_SIZE
 - ATTR_PM_SPIPSS_IN_DELAY
 - ATTR_PM_SPIPSS_INTER_FRAME_DELAY
 - Output feature attributes
 - ATTR_PM_SPIPSS_CLOCK_DIVIDER
 - ATTR_PM_SPIPSS_OUT_COUNT
 - ATTR_PM_SPIPSS_IN_COUNT
 - ATTR_PM_SPIPSS_INTER_FRAME_DELAY_SETTING
- b) p8_pm_init.C *chiptarget, ENUM:PM_INIT*
- *p8_pcbs_init.C *chiptarget, ENUM:PM_INIT*
 - Resets DPLL_CPM_PARM_REG.cpm_filter_enable
 - Sleep configuration
 - ATTR_PM_SLEEP_TYPE (Deep or Fast)
 - ATTR_PM_SLEEP_ENTRY (Assisted or Hardware) – depends on di/dt characteristics of the system (Assisted if power off serialization is needed, Hardware if the system can handle the unrelated powering off between cores. Hardware decreases entry latency)
 - ATTR_PM_SLEEP_EXIT (Assisted or Hardware) – set to Assisted (for both Fast and Deep). Fast for di/dt management; Deep as this necessary for restore. Setting to Hardware is a test mode for Fast only.
 - Winkle configuration
 - ATTR_PM_WINKLE_TYPE (Deep or Fast)
 - ATTR_PM_WINKLE_ENTRY (Assisted or Hardware)
 - ATTR_PM_WINKLE_EXIT (Assisted or Hardware) – set to Assisted (for both Fast and Deep). Fast for di/dt management; Deep as this necessary for restore. Setting to Hardware is a test mode for Fast only.

- PMCR default value adjustment (Hardware flush 0 -> [restore to 0 for reset case](#))
 - For reset case, disable all “global_en” bits in PMCR and PMICR; this keeps Global Pstate Request from occurring to the PMC until it has been initialized. OCCFW to be do this.
- PMICR default value adjustment (Hardware flush 0 -> restore to TBD for reset) **TODO**
 - How does policy influence the PMICR Pstate values?
 - Base: run at the turbo value fixed
 - Enhancement: run at the highest Pstate value on the chip. (needs power projection to judge worth).
 - latency enable
 - Not planned at this time.
- *p8roc_pcbs_init_resclk.C*
 - Overlay any stepping program changes from the flush value
 - Step delays (MRWB attribute)
 - Leaves resonant clocking function disabled.
- *p8_pcbs_ivrm.C* (If enabled) ([OCC Config Data area item: ivrms enabled/disabled](#))
 - *p8roc_pcbs_ivrm_cal.C*
 - Calibration of PVREF (procedure) ([FAPI](#))
 - Note: this is done via FAPI (at IPL) as the calibration circuits are on VIO in the EX chiplet and thus are NOT part of the Winkle image.
 - Note: if a core is “brought online” after IPL and the iVRMs are going to be used, this procedure must be run targeting that chiplet
- *p8_pmc_init.C *chiptarget, ENUM:PM_INIT*
 - PMC Hang Pulse setup (based on nest frequency) (MRWB Attribute: Nest Frequency MHz)
 - PMC PORE-SW Timeout setup (based on nest frequency)
 - PMC SPIVID Configuration
 - Number of links ([ATTR_PM_SPIVID_PORT_ENABLE](#))
 - Write Status Delay ([ATTR_PM_SPIVID_INTERFRAME_DELAY_WRITE_STATUS_VALUE](#))
 - Read State Delay ([MRWB attribute](#))
- *p8_oha_init.C *chiptarget, ENUM:PM_INIT*
 - OHA Idle timeout settings
 - **ATTR_PM_AISS_TIMEOUT - ENUM:** 1ms=0, 2ms=1, 4ms=2, 8ms=3, 16ms=4, 32ms=5, 64ms=6, 128ms=7, 256ms=8, 512ms=9; other values are illegal
 - Default: 32ms to cover the time of a Deep Sleep (5ms) plus a Deep Winkle (10ms) x 2.
 - Power Proxy Trace Time Select (MWRB Attribute) 0=0.25us , 1=0.5us, 2=1us, and 3=2us
 - Default: 1us
- *p8_pss_init.C *chiptarget, ENUM:PM_INIT*
 - PSS Configuration
 - **ATTR_PM_SPIPSS_CLOCK_DIVIDER**

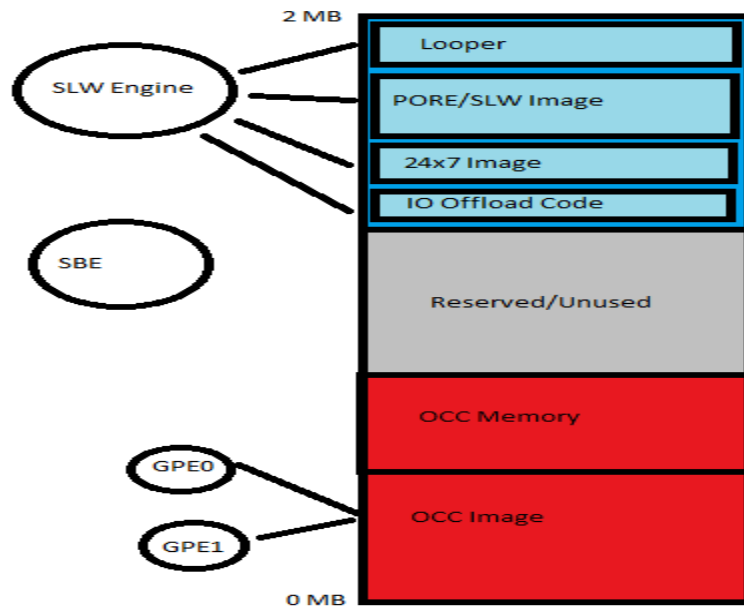
- ATTR_PM_SPIPSS_OUT_COUNT
- ATTR_PM_SPIPSS_IN_COUNT
- ATTR_PM_SPIPSS_INTER_FRAME_DELAY_SETTING
- *p8_pba_init.C *chiptarget, ENUM:PM_INIT*
 - “PowerBus Slave” buffer set configuration. Assigns slaves to OCI masters for runtime (vs IPL time for HBI loading)
 - PBA Configuration
 - Hang pulse dividers
 - Drop priority (MRWB attribute - TBD)
 - Overcommit counter settings (MRWB attribute - TBD)
 - Setup PBA for unicast/multicast
 - ATTR_PM_PBAX_NODEID
 - ATTR_PM_PBAX_CHIPID
 - ATTR_PM_PBAX_BRDCST_GROUPID
- *p8_occ_init.C *chiptarget, ENUM:OCC_CONFIG*
 - Place holder only at this point
 - OCCFW takes care of:
 - OCB Interrupt clearing
 - Indirect Channel reset
- *p8_pm_firinit.C* : Set the FIR masks and action bits per RAS FIR spreadsheet; done as FAPIs vs scom.initfiles to be supportable under PHYP
 - *p8_pm_pmc_firinit.C *chiptarget*
 - *p8_pm_pcbs_firinit.C *chiptarget*
 - *p8_pm_oha_firinit.C *chiptarget*
 - *p8_pm_occ_firinit.C *chiptarget*

24. Load OCC

24.1. load_occ : Place OCC image into memory

For each chip in a physical node

- There are two divergent paths to load the OCC code image. The first is lab/Cronus only without FW. In this case the HWP is run. In the second case FW controls building up the image at the direction of PHYP
 - *p8roc_occ_load.C* **CRONUS**
 - Load image in memory from PNOR at an address that is passed to this procedure
 - *occ_load:* **FW**
 - PHYP will call host services with memory region to place the HOMER image
 - Host services – request OCC, SLW images from FSP via lidmgr
 - PHYP will verify signature through secure algorithm
 - Host Services will customize SLW image (see step 15 of IPL)
 - Host Services will place OCC initial startup information into HOMER image
 - Nest Frequency
 - Host services place SLW and OCC images as directed by PHYP. Here is an overview of a completed imagewill layout HOMER image like this:



25. Start OCC

25.1. `start_occ` : Start OCC

- *`p8roc_occ_control.C *chiptarget, ENUM:OCC_START`*
 - Starts OCC load. by releasing the reset to the PPC405.
 - For DCMs, the Slave OCC must be released first;
 - Slave/Master
 - Put PS data to Slave
 - OCC in Slave consume the data and init hardware (load LPST, init PMC, init resclk...)
 - Poll (attn) that OCC is done with install
 - Master
 - OCC code bootloads itself from Memory into SRAM tank
 - Also need
 - *`proc_occ_control.C *chiptarget, ENUM:OCC_STOP`*
 - *`p8_pm_heartbeat.C`*

26. Config OCC

26.1. `config_occ` : Load OCC config

TMGT (FSP or Host) now builds the OCC config data and uses its communication path to OCC to load

a) *`p8_build_pstate_datablock.C *chiptarget, *pstate_superstructure`*

- Manufacturing request to allow biasing
- **`build_gpstate_table()`**
 - Read VPD attributes and compute the Global PState Table as a data block. Will be installed in OCC SRAM by OCCFW.
 - Input platform attributes
 - **`ATTR_FREQ_CORE_MAX`**

- **ATTR_PROC_R_LOADLINE**
- **ATTR_PROC_R_DISTLOSS**
- **ATTR_PROC_VRM_VOFFSET**
- **ATTR_MVPD_VOLTAGE_CONTROL** (array from #V)
- **ATTR_FREQ_PROC_REFCLOCK** (*freq_attributes.xml*)
- **ATTR_DPLL_DIVIDER**
 - Forward divide * feedback divide – typically 8
- External Voltage Bias
 - % from nominal up (**ATTR_VOLTAGE_EXT_BIAS_UP**)
 - % from nominal down (**ATTR_VOLTAGE_EXT_BIAS_DOWN**)
 - If both are specified, return an error (due to unsigned attributes)
- Internal Voltage Bias
 - % from nominal up (**ATTR_VOLTAGE_INT_BIAS_UP**)
 - % from nominal down (**ATTR_VOLTAGE_INT_BIAS_DOWN**)
 - If both are specified, return an error (due to unsigned attributes)
- Frequency
 - % from nominal up (**ATTR_FREQ_BIAS_UP**)
 - % from nominal down (**ATTR_FREQ_BIAS_DOWN**)
 - If both are specified, return an error (due to unsigned attributes)
- Output feature attributes
 - **ATTR_PM_PSTATE0_FREQUENCY**
 - Output Configuration Data (Pstate Superstructure)
 - **PSTATE0_FREQUENCY_KHZ** (4B)
 - **PSTATE0_FREQUENCY_CODE** (1B array x 16 cores)
 - **DPLL_FMAX_BIAS** (1B array x 16 cores)
 - Parm to procedure to allow characterization to bias individual cores for CPMs
 - Number of entries (1B)
 - Frequency step in KHz (4B)
 - Min Pstate in GPST (1B)
 - Undervolting bias (1B)
 - **Global Pstate Table** (128 x 8B = 1024B)
- **build_lpst_table()**
 - Build the Local Pstate Table and VDS table (hardware array per chiplet loaded with common data) based on VPD #M attributes. Will be passed to OCCFW as part of the OCC config data for installation in the hardware by an OCC applet.
 - Uses the Least Squares fit algorithm from Josh for VDS table
 - Based on the maximum supported frequency (**ATTR_FREQ_CORE_MAX**), the refclk frequency (**ATTR_FREQ_PROC_REFCLOCK** and **DPLL divider** (8) (**TBD: check on attr**), compute *lpsi_min* to place the LPST into Pstate space. Set *lpsi_entries_minus_1* to 127.

- Input platform attributes
 - **ATTR_FREQ_CORE_MAX**
 - **ATTR_FREQ_PROC_REFCLOCK**
 - **ATTR_DPLL_DIVIDER**
 - **ATTR_MVPD_IVRM_CALIBRATION** (array from #M)
- Configuration Data
 - **LPSI_MIN** (1B)
 - **LPSI_ENTRIES** (1B)
 - **lpst_load()** OCC applet must adjust this to subtract 1 before loading into the hardware
 - **Local Pstate Table (physical table)** (96 x 8B = 768B)
 - LSPT
 - VDS
 - Multiplier table
- **convert_safe_freq()**
 - PState translation of safe frequency using **PSTATE0_FREQUENCY**
 - Input platform attributes
 - **ATTR_PM_SAFE_FREQUENCY** (binary in MHz)
 - **ATTR_PM_PSTATE0_FREQUENCY** (binary in Mhz)
 - Configuration Data (1B)
 - **SAFE_PSTATE** (1B)
- **convert_safe_voltage()**
 - Set the PMC Vsafe value
 - Input platform attributes
 - **ATTR_PM_SAFE_VOLTAGE**
 - Configuration Data (1B)
 - **PVSAFE_PSTATE** (1B)
- **convert_resclk_freqs()**
 - Convert the Resonant Clocking band definitions from frequency to Pstates. The conversion uses **ATTR_PM_PSTATE0_FREQUENCY**.
 - Input platform attributes
 - **ATTR_PM_RESONANT_CLOCK_FULL_CLOCK_SECTOR_BUFFER_FREQUENCY**
 - **ATTR_PM_RESONANT_CLOCK_LOW_BAND_LOWER_FREQUENCY**
 - **ATTR_PM_RESONANT_CLOCK_LOW_BAND_UPPER_FREQUENCY**
 - **ATTR_PM_RESONANT_CLOCK_HIGH_BAND_LOWER_FREQUENCY**
 - **ATTR_PM_RESONANT_CLOCK_HIGH_BAND_UPPER_FREQUENCY**
 - Configuration Data (5B)
 - **RESONANT_CLOCK_FULL_CSB_PSTATE** (8B)
 - **RESONANT_CLOCK_LFRLOW_PSTATE** (8B)
 - **RESONANT_CLOCK_LFRUPPER_PSTATE** (8B)
 - **RESONANT_CLOCK_HFRLOW_PSTATE** (8B)

- `RESONANT_CLOCK_HFRHIGH_PSTATE` (8B)
- **build_cpm_data()**
 - Read CPM (#P) attributes and convert these to SCOM calibration settings stored in the superstructure.
- Build OCC config data block
 - Good cores come via the deconfig register.